

# Reading 23 Assignment (Nov 24)

CSE 291A

John Lee

Summary:

Amazon Bedrock AgentCore is a platform for operating and deploying AI agents at scale, which includes a Python SDK and toolkit that enables developers to build and deploy AI agents. The platform works with LangGraph as well as custom frameworks with zero need for existing servers or containers, allowing it to be production-ready from the get-go. AgentCore also provides many enterprise-grade services such as persistent knowledge, secure cloud-based browser, secure auth, and serverless deployment.

1. Can you think of some new challenges or opportunities when serving agents in a serverless way from traditional serverless services like AWS Lambda?
  - o The main difference in AWS Lambda serverless service is that it is stateless and does not have persistent memory like Bedrock AgentCore. AgentCore being able to have built-in state management is a great opportunity. AWS Lambda having strict execution times is a challenge for agent serving.
2. What do you think are the key requirements for agent hosting?
  - o As mentioned in the summary, AgentCore offers many enterprise-grade services such as persistent knowledge, secure auth, and real-time monitoring. These factors such as state persistence (for maintaining conversation continuity), memory, and secure authentication are key requirements for agent hosting.